

Hands-on Activity 4.2 Transfer Learning on PyTorch	
Course Code: CPE 313	Program: BSCPE
Course Title: Advance Machine Learning and Deep Learning	Date Performed: 02/19/2026
Section: CPE32S3	Date Submitted: 02/19/2026
Name: REUEL CHRISTIAN PORNObE	Instructor: Engr. NEAL BARTON JAMES MATIRA
1. Discussion	
Objectives: This activity aims to demonstrate the use of PyTorch in transfer learning to solve a given problem.	
2. Materials and Equipment	
<ul style="list-style-type: none"> - Computer - Internet 	
3. Procedure	
Follow the instructions as provided in the notebook attached in the instruction materials. Output: Provide a lab report showing: All procedures followed (with screenshots) in the notebookLinks to an external site.. Supplementary Activity: Answer to the following supplementary activity: Choose a pretrained model. Finetune on your dataset from the previous activity. Evaluate the performance of the previous model to this finetuned model. Utilize the pretrained ConvNet model as fixed feature extractor. Evaluate the performance of the previous model to this finetuned model. Discuss the following: How did finetuning affect your performance? Which of the different situations for rule of thumb were applicable to you?	
4. Output	

```
[1] ✓ 0s
    ▾ Transfer Learning using PyTorch
      ↴
      1 !nvidia-smi
      Thu Feb 19 05:55:14 2026
      +-----+
      | NVIDIA-SMI 580.82.07      Driver Version: 580.82.07      CUDA Version: 13.0 |
      +-----+
      | GPU  Name        Persistence-M  Bus-Id      Disp.A  Volatile Uncorr. ECC  |
      | Fan  Temp  Perf  Pwr:Usage/Cap | Memory-Usage | GPU-Util  Compute M. |
      |          |          |          |          |      MIG M.   |
      +-----+
      | 0  Tesla T4           Off  00000000:00:04.0 Off   0MiB / 15360MiB   0%  Default |
      | N/A  52C   P8          11W / 70W |                  |          |          |
      +-----+
      Processes:
      GPU  GI  CI          PID  Type  Process name          GPU Memory Usage
      ID  ID
      +-----+
      | no running processes found
      +-----+



[2] ✓ 9s
    ▾ Load Important Libraries
      ↴
      1 # License: BSD
      2 # Author: Sasank Chilamkurthy
      3
      4 import torch
      5 import torch.nn as nn
      6 import torch.optim as optim
      7 from torch.optim import lr_scheduler
      8 import torch.backends.cudnn as cudnn
      9 import numpy as np
     10 import torchvision
     11 from torchvision import datasets, models, transforms
     12 import matplotlib.pyplot as plt
     13 import time
     14 import os
     15 from PIL import Image
     16 from tempfile import TemporaryDirectory
     17
     18 cudnn.benchmark = True
     19 plt.ion()  # interactive mode
      ↴
      <contextlib.ExitStack at 0x7e0de0ff08c0>
```

Importing all the necessary Libraries

Load Data

=====

We will use torchvision and torch.utils.data packages for loading the data.

The problem we're going to solve today is to train a model to classify **ants** and **bees**. We have about 120 training images each for ants and bees. There are 75 validation images for each class. Usually, this is a very small dataset to generalize upon, if trained from scratch. Since we are using transfer learning, we should be able to generalize reasonably well.

This dataset is a very small subset of imagenet.

```
[3]  ✓ 0s   1 rm -R /content/hymenoptera_data/train/.ipynb_checkpoints
2 ls /content/hymenoptera_data/test/train -a #to make sure that the deletion has occurred
3
4 rm -R /content/hymenoptera_data/val/.ipynb_checkpoints
5 ls /content/hymenoptera_data/val -a #to make sure that the deletion has occurred
...
... rm: cannot remove '/content/hymenoptera_data/train/.ipynb_checkpoints': No such file or directory
ls: cannot access '/content/hymenoptera_data/test/train': No such file or directory
rm: cannot remove '/content/hymenoptera_data/val/.ipynb_checkpoints': No such file or directory
ls: cannot access '/content/hymenoptera_data/val': No such file or directory
[4]  ✓ 2s   1 from google.colab import drive
2 drive.mount('/content/drive')
...
Mounted at /content/drive
[5]  ✓ 2s   1 # Data augmentation and normalization for training
2 # Just normalization for validation
3 data_transforms = {
4     'train': transforms.Compose([
5         transforms.RandomResizedCrop(224),
6         transforms.RandomHorizontalFlip(),
7         transforms.ToTensor(),
8         transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
9     ]),
10    'val': transforms.Compose([
11        transforms.Resize(256),
12        transforms.CenterCrop(224),
13        transforms.ToTensor(),
14        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
15    ]),
16 }
17
18 data_dir = '/content/drive/MyDrive/hymenoptera_data/hymenoptera_data'
19 image_datasets = {x: datasets.ImageFolder(os.path.join(data_dir, x),
20                                         data_transforms[x])
21     for x in ['train', 'val']}
22 dataloaders = {x: torch.utils.data.DataLoader(image_datasets[x], batch_size=4,
23                                               shuffle=True, num_workers=4)
24     for x in ['train', 'val']}
25 dataset_sizes = {x: len(image_datasets[x]) for x in ['train', 'val']}
26 class_names = image_datasets['train'].classes
27
```

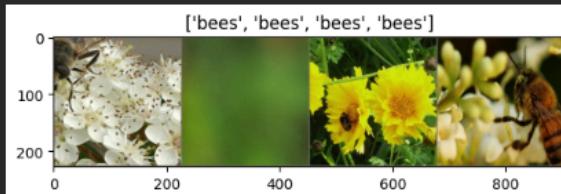
Loading and transforming the data for Resnet

▼ Visualize a few images

=====

Let's visualize a few training images so as to understand the data augmentations.

```
[7]  ✓ 0s
  1 def imshow(inp, title=None):
  2     """Display image for Tensor."""
  3     inp = inp.numpy().transpose((1, 2, 0))
  4     mean = np.array([0.485, 0.456, 0.406])
  5     std = np.array([0.229, 0.224, 0.225])
  6     inp = std * inp + mean
  7     inp = np.clip(inp, 0, 1)
  8     plt.imshow(inp)
  9     if title is not None:
 10         plt.title(title)
 11     plt.pause(0.001) # pause a bit so that plots are updated
 12
 13
 14 # Get a batch of training data
 15 inputs, classes = next(iter(dataloaders['train']))
 16
 17 # Make a grid from batch
 18 out = torchvision.utils.make_grid(inputs)
 19
 20 imshow(out, title=[class_names[x] for x in classes])
```



Created a function to visualize images from the data

▼ Visualizing the model predictions

=====

Generic function to display predictions for a few images

```
[8]  ✓ 0s
  1 def visualize_model(model, num_images=6):
  2     was_training = model.training
  3     model.eval()
  4     images_so_far = 0
  5     fig = plt.figure()
  6
  7     with torch.no_grad():
  8         for i, (inputs, labels) in enumerate(dataloaders['val']):
  9             inputs = inputs.to(device)
 10             labels = labels.to(device)
 11
 12             outputs = model(inputs)
 13             _, preds = torch.max(outputs, 1)
 14
 15             for j in range(inputs.size()[0]):
 16                 images_so_far += 1
 17                 ax = plt.subplot(num_images//2, 2, images_so_far)
 18                 ax.axis('off')
 19                 ax.set_title(f'predicted: {class_names[preds[j]]}')
 20                 imshow(inputs.cpu().data[j])
 21
 22             if images_so_far == num_images:
 23                 model.train(mode=was_training)
 24                 return
 25
  model.train(mode=was_training)
```

Visualize the model prediction and image

▼ Training the model

```
=====
Now, let's write a general function to train a model. Here, we will illustrate:
  • Scheduling the learning rate
  • Saving the best model
In the following, parameter scheduler is an LR scheduler object from torch.optim.lr_scheduler.
```

```
[8]  ✓ 0s
1 def train_model(model, criterion, optimizer, scheduler, num_epochs=25):
2     since = time.time()
3
4     # Create a temporary directory to save training checkpoints
5     with TemporaryDirectory() as tempdir:
6         best_model_params_path = os.path.join(tempdir, 'best_model_params.pt')
7
8         torch.save(model.state_dict(), best_model_params_path)
9         best_acc = 0.0
10
11        for epoch in range(num_epochs):
12            print(f'Epoch {epoch}/{num_epochs - 1}')
13            print('-' * 10)
14
15            # Each epoch has a training and validation phase
16            for phase in ['train', 'val']:
17                if phase == 'train':
18                    model.train() # Set model to training mode
19                else:
20                    model.eval() # Set model to evaluate mode
21
22                running_loss = 0.0
23                running_corrects = 0
24
25                # Iterate over data.
26                for inputs, labels in dataloaders[phase]:
27                    inputs = inputs.to(device)
28                    labels = labels.to(device)
29
30                    # zero the parameter gradients
31                    optimizer.zero_grad()
32
33                    # forward
34                    # track history if only in train
35                    with torch.set_grad_enabled(phase == 'train'):
36                        outputs = model(inputs)
37                        _, preds = torch.max(outputs, 1)
38                        loss = criterion(outputs, labels)
39
40
41    # backward + optimize only if in training phase
42    loss.backward()
43    optimizer.step()
44
45    # track progress
46    if phase == 'train':
47        scheduler.step()
48
49    else:
50        acc = torchmetrics.functional.accuracy(preds, labels)
51
52        if acc > best_acc:
53            best_acc = acc
54            torch.save(model.state_dict(), best_model_params_path)
55
56    print(f'{phase} Loss: {running_loss / len(dataloaders[phase]):.4f} | Accuracy: {best_acc:.4f}%')
57
58 print(f'Training complete in {time.time() - since} mins')
59 print(f'Best accuracy: {best_acc:.4f}%')
```

Creates a function to train a model

▼ Visualizing the model predictions

=====

Generic function to display predictions for a few images

```
[9] ✓ 0s 1 def visualize_model(model, num_images=6):
2     was_training = model.training
3     model.eval()
4     images_so_far = 0
5     fig = plt.figure()
6
7     with torch.no_grad():
8         for i, (inputs, labels) in enumerate(dataloaders['val']):
9             inputs = inputs.to(device)
10            labels = labels.to(device)
11
12            outputs = model(inputs)
13            _, preds = torch.max(outputs, 1)
14
15            for j in range(inputs.size()[0]):
16                images_so_far += 1
17                ax = plt.subplot(num_images//2, 2, images_so_far)
18                ax.axis('off')
19                ax.set_title(f'predicted: {class_names[preds[j]]}')
20                imshow(inputs.cpu().data[j])
21
22            if images_so_far == num_images:
23                model.train(mode=was_training)
24                return
25
model.train(mode=was_training)
```

▼ Finetuning the ConvNet

=====

Load a pretrained model and reset final fully connected layer.

```
[10] ✓ 0s 1 model_ft = models.resnet18(weights='IMAGENET1K_V1')
2 num_ftrs = model_ft.fc.in_features
3 # Here the size of each output sample is set to 2.
4 # Alternatively, it can be generalized to ``nn.Linear(num_ftrs, len(class_names))``.
5 model_ft.fc = nn.Linear(num_ftrs, 2)
6
7 model_ft = model_ft.to(device)
8
9 criterion = nn.CrossEntropyLoss()
10
11 # Observe that all parameters are being optimized
```

Fine tune the Resnet model for the hymenoptera data

Train and evaluate

It should take around 15-25 min on CPU. On GPU though, it takes less than a minute.

```
[11] 1 model_ft = train_model(model_ft, criterion, optimizer_ft, exp_lr_scheduler,  
2                               num_epochs=25)
```

... Show hidden output

```
[12] 1 visualize_model(model_ft)
```



Model training of the Fine tuned model

ConvNet as fixed feature extractor

Here, we need to freeze all the network except the final layer. We need to set `requires_grad = False` to freeze the parameters so that the gradients are not computed in `backward()`.

You can read more about this in the documentation [here](#).

```
[13] 1 model_conv = torchvision.models.resnet18(weights='IMAGENET1K_V1')  
2 for param in model_conv.parameters():  
3     param.requires_grad = False  
4  
5 # Parameters of newly constructed modules have requires_grad=True by default  
6 num_ftrs = model_conv.fc.in_features  
7 model_conv.fc = nn.Linear(num_ftrs, 2)  
8  
9 model_conv = model_conv.to(device)  
10  
11 criterion = nn.CrossEntropyLoss()  
12  
13 # Observe that only parameters of final layer are being optimized as  
14 # opposed to before.  
15 optimizer_conv = optim.SGD(model_conv.fc.parameters(), lr=0.001, momentum=0.9)  
16  
17 # Decay LR by a factor of 0.1 every 7 epochs  
18 exp_lr_scheduler = lr_scheduler.StepLR(optimizer_conv, step_size=7, gamma=0.1)
```

Trains only the last layer of the Resnet Model

```
Choose a pretrained model.

1 from torchvision.models import inception_v3
2 supp_model = torchvision.models.inception_v3(weights='IMAGENET1K_V1')

3 # Data for Inception V3
4 data_transforms = {
5     'train': transforms.Compose([
6         transforms.RandomResizedCrop(299),
7         transforms.RandomHorizontalFlip(),
8         transforms.ToTensor(),
9         transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
10    ]),
11    'val': transforms.Compose([
12        transforms.Resize(299),
13        transforms.CenterCrop(299),
14        transforms.ToTensor(),
15        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
16    ]),
17 data_dir = '/content/drive/MyDrive/hymenoptera_data/hymenoptera_data'
18 image_datasets = {x: datasets.ImageFolder(os.path.join(data_dir, x),
19                                         data_transforms[x])
20     for x in ['train', 'val']}
21 dataloaders = {x: torch.utils.data.DataLoader(image_datasets[x], batch_size=4,
22                                                 shuffle=True, num_workers=4)
23     for x in ['train', 'val']}
24 dataset_sizes = {x: len(image_datasets[x]) for x in ['train', 'val']}
25 class_names = image_datasets['train'].classes
26
27 # We want to be able to train our model on an `accelerator <https://pytorch.org/docs/stable/torch.html#accelerators>`_
28 # such as CUDA, MPS, MTIA, or XPU. If the current accelerator is available, we will use it. Otherwise, we use the CPU.
29
30 device = torch.accelerator.current_accelerator().type if torch.accelerator.is_available() else "cpu"
31 print(f"Using {device} device")
32
33 ... Using cuda device
```

In supplementary Activity, I chose Inception_V3 then changed the Data_transform from 256 -> 299.

Finetune on your dataset from the previous activity.

```
[?2] ✓ 0s
  1 import torchvision
  2 import torch.nn as nn
  3 import torch.optim as optim
  4 from torch.optim import lr_scheduler
  5
  6 supp_model_ft = torchvision.models.inception_v3(
  7     weights='IMAGENET1K_V1',
  8     aux_logits=True
  9 )
 10
 11 # Replace main classifier
 12 num_ftrs = supp_model_ft.fc.in_features
 13 supp_model_ft.fc = nn.Linear(num_ftrs, 2)
 14
 15 # Replace auxiliary classifier
 16 num_ftrs_aux = supp_model_ft.AuxLogits.fc.in_features
 17 supp_model_ft.AuxLogits.fc = nn.Linear(num_ftrs_aux, 2)
 18
 19 supp_model_ft = supp_model_ft.to(device)
 20
 21 criterion = nn.CrossEntropyLoss()
 22
 23 optimizer_ft = optim.SGD(
 24     supp_model_ft.parameters(),
 25     lr=0.001,
 26     momentum=0.9
 27 )
 28
 29 exp_lr_scheduler = lr_scheduler.StepLR(
 30     optimizer_ft,
 31     step_size=7,
 32     gamma=0.1
 33 )
 34
```

Fine tune for Inception V3

Evaluate the performance of the previous model to this finetuned model.

```
[75] ✓ 0s
  1 def train_model_inception_v3(model, criterion, optimizer, scheduler, num_epochs=25):
  2     since = time.time()
  3
  4     with TemporaryDirectory() as tempdir:
  5         best_model_params_path = os.path.join(tempdir, 'best_model_params.pt')
  6         torch.save(model.state_dict(), best_model_params_path)
  7
  8         best_acc = 0.0
  9
 10        for epoch in range(num_epochs):
 11            print(f'Epoch {epoch}/{num_epochs - 1}')
 12            print('-' * 10)
 13
 14            for phase in ['train', 'val']:
 15                if phase == 'train':
 16                    model.train()
 17                else:
 18                    model.eval()
 19
 20                running_loss = 0.0
 21                running_corrects = 0
 22
 23                for inputs, labels indataloaders[phase]:
 24                    inputs = inputs.to(device)
 25                    labels = labels.to(device)
 26
 27                    optimizer.zero_grad()
 28
 29                    with torch.set_grad_enabled(phase == 'train'):
 30                        outputs = model(inputs)
 31
 32                        if phase == 'train' and hasattr(outputs, 'aux_logits'):
 33                            loss1 = criterion(outputs.logits, labels)
 34                            loss2 = criterion(outputs.aux_logits, labels)
 35                            loss = loss1 + 0.4 * loss2
 36                            logits = outputs.logits
 37
 38                        else:
 39                            if hasattr(outputs, 'logits'):
 40                                logits = outputs.logits
 41                            else:
 42                                logits = outputs
 43                            loss = criterion(logits, labels)
 44
 45                            _, preds = torch.max(logits, 1)
 46
 47                            if phase == 'train':
 48                                loss.backward()
 49                                optimizer.step()
 50
 51                            running_loss += loss.item() * inputs.size(0)
```

```
[76] ✓ 3m
  1 supp_model_ft = train_model_inception_v3(supp_model_ft, criterion, optimizer_ft, exp_lr_scheduler,
  2                                         num_epochs=25)
```

> ... Show hidden output

Train Inception V3 with a new train function for it.

Utilize the pretrained ConvNet model as fixed feature extractor.

```

[79]   1 import torchvision
  2 import torch.nn as nn
  3 import torch.optim as optim
  4 from torch.optim import lr_scheduler
  5
  6 model_conv = torchvision.models.inception_v3(
  7     weights='IMAGENET1K_V1',
  8     aux_logits=True
  9 )
10
11 # Freeze all layers
12 for param in model_conv.parameters():
13     param.requires_grad = False
14
15 # Replace main classifier
16 num_ftrs = model_conv.fc.in_features
17 model_conv.fc = nn.Linear(num_ftrs, 2)
18
19 # Replace auxiliary classifier
20 num_ftrs_aux = model_conv.AuxLogits.fc.in_features
21 model_conv.AuxLogits.fc = nn.Linear(num_ftrs_aux, 2)
22
23 model_conv = model_conv.to(device)
24
25 criterion = nn.CrossEntropyLoss()
26
27 # Only optimize the new classifier layers
28 optimizer_conv = optim.SGD(
29     list(model_conv.fc.parameters()) + list(model_conv.AuxLogits.fc.parameters()),
30     lr=0.001,
31     momentum=0.9
32 )
33
34 exp_lr_scheduler = lr_scheduler.StepLR(
35     optimizer_conv,
36     step_size=7,
37     gamma=0.1
38 )
39
```

[80] 1 model_conv = train_model_inception_v3(model_conv, criterion, optimizer_conv,
 2 exp_lr_scheduler, num_epochs=25)

Trains the Inception V3 with only its last layer.

Procedure Model 1
Training complete in 2m 32s Best val Acc: 0.941176
Procedure Model 2
Val Loss: 0.1655 Acc: 0.9477 Training complete in 1m 40s Best val Acc: 0.960784
Supplementary Model 1
Training complete in 3m 52s Best val Acc: 0.9542
Supplementary Model 2

Training complete in 2m 41s
Best val Acc: 0.9412

5. Supplementary Activity

In supplementary Activity, I chose Inception_V3 then changed the Data_transform from 254 -> 299.

Finetune on your dataset from the previous activity.

```
[72] ✓ 0s
  1 import torchvision
  2 import torch.nn as nn
  3 import torch.optim as optim
  4 from torch.optim import lr_scheduler
  5
  6 supp_model_ft = torchvision.models.inception_v3(
  7     weights='IMAGENET1K_V1',
  8     aux_logits=True
  9 )
 10
 11 # Replace main classifier
 12 num_ftrs = supp_model_ft.fc.in_features
 13 supp_model_ft.fc = nn.Linear(num_ftrs, 2)
 14
 15 # Replace auxiliary classifier
 16 num_ftrs_aux = supp_model_ft.AuxLogits.fc.in_features
 17 supp_model_ft.AuxLogits.fc = nn.Linear(num_ftrs_aux, 2)
 18
 19 supp_model_ft = supp_model_ft.to(device)
 20
 21 criterion = nn.CrossEntropyLoss()
 22
 23 optimizer_ft = optim.SGD(
 24     supp_model_ft.parameters(),
 25     lr=0.001,
 26     momentum=0.9
 27 )
 28
 29 exp_lr_scheduler = lr_scheduler.StepLR(
 30     optimizer_ft,
 31     step_size=7,
 32     gamma=0.1
 33 )
 34
```

Fine tune for Inception V3

Evaluate the performance of the previous model to this finetuned model.

```
[75] ✓ 0s
  1 def train_model_inception_v3(model, criterion, optimizer, scheduler, num_epochs=25):
  2     since = time.time()
  3
  4     with TemporaryDirectory() as tempdir:
  5         best_model_params_path = os.path.join(tempdir, 'best_model_params.pt')
  6         torch.save(model.state_dict(), best_model_params_path)
  7
  8         best_acc = 0.0
  9
 10        for epoch in range(num_epochs):
 11            print(f'Epoch {epoch}/{num_epochs - 1}')
 12            print('-' * 10)
 13
 14            for phase in ['train', 'val']:
 15                if phase == 'train':
 16                    model.train()
 17                else:
 18                    model.eval()
 19
 20                running_loss = 0.0
 21                running_corrects = 0
 22
 23                for inputs, labels indataloaders[phase]:
 24                    inputs = inputs.to(device)
 25                    labels = labels.to(device)
 26
 27                    optimizer.zero_grad()
 28
 29                    with torch.set_grad_enabled(phase == 'train'):
 30                        outputs = model(inputs)
 31
 32                        if phase == 'train' and hasattr(outputs, 'aux_logits'):
 33                            loss1 = criterion(outputs.logits, labels)
 34                            loss2 = criterion(outputs.aux_logits, labels)
 35                            loss = loss1 + 0.4 * loss2
 36                            logits = outputs.logits
 37
 38                        else:
 39                            if hasattr(outputs, 'logits'):
 40                                logits = outputs.logits
 41                            else:
 42                                logits = outputs
 43                            loss = criterion(logits, labels)
 44
 45                        _, preds = torch.max(logits, 1)
 46
 47                        if phase == 'train':
 48                            loss.backward()
 49                            optimizer.step()
 50
 51                running_loss += loss.item() * inputs.size(0)
```

```
[76] ✓ 3m
  1 supp_model_ft = train_model_inception_v3(supp_model_ft, criterion, optimizer_ft, exp_lr_scheduler,
  2                                         num_epochs=25)
```

> ... Show hidden output

Train Inception V3 with a new train function for it.

Utilize the pretrained ConvNet model as fixed feature extractor.

```

[79]   1 import torchvision
  2 import torch.nn as nn
  3 import torch.optim as optim
  4 from torch.optim import lr_scheduler
  5
  6 model_conv = torchvision.models.inception_v3(
  7     weights='IMAGENET1K_V1',
  8     aux_logits=True
  9 )
10
11 # Freeze all layers
12 for param in model_conv.parameters():
13     param.requires_grad = False
14
15 # Replace main classifier
16 num_ftrs = model_conv.fc.in_features
17 model_conv.fc = nn.Linear(num_ftrs, 2)
18
19 # Replace auxiliary classifier
20 num_ftrs_aux = model_conv.AuxLogits.fc.in_features
21 model_conv.AuxLogits.fc = nn.Linear(num_ftrs_aux, 2)
22
23 model_conv = model_conv.to(device)
24
25 criterion = nn.CrossEntropyLoss()
26
27 # Only optimize the new classifier layers
28 optimizer_conv = optim.SGD(
29     list(model_conv.fc.parameters()) + list(model_conv.AuxLogits.fc.parameters()),
30     lr=0.001,
31     momentum=0.9
32 )
33
34 exp_lr_scheduler = lr_scheduler.StepLR(
35     optimizer_conv,
36     step_size=7,
37     gamma=0.1
38 )
39
```

[80] 1 model_conv = train_model_inception_v3(model_conv, criterion, optimizer_conv,
 2 exp_lr_scheduler, num_epochs=25)

Trains the Inception V3 with only its last layer.

Supplementary Model 1

Training complete in 3m 52s
 Best val Acc: 0.9542

Supplementary Model 2

Training complete in 2m 41s
 Best val Acc: 0.9412

6. Conclusion

In this lab activity, I did transfer learning of the Resnet for classifying the dataset, hymenoptera. I learned how it works first hand and how to only change or finetune the last layers of the existing model. For the supplementary activity, I used inception v3 which has different transform size compared to the Resnet model. Therefore, I changed the sizes of data from 254 -> 299. I also

changed how the train_model function works for inception v3. Overall, I was able to understand how to do basic Transfer learning of data.

6. Assessment Rubric

Lab Activity Rubric							
Criteria	Ratings						Pts
⌚ SO 7 PI 1 Student Outcome 7.1 Acquire and apply new knowledge from outside sources. threshold: 4.8 pts	6 pts Excellent Educational interests and pursuits exist and flourish outside classroom requirements,knowledge and/or experiences are pursued independently and applies knowledge learned into practice	5 pts Good Educational interests and pursuits exist and flourish outside classroom requirements,knowledge and/or experiences are pursued independently	4 pts Satisfactory Look beyond classroom requirements, showing interest in pursuing knowledge independently	3 pts Unsatisfactory Begins to look beyond classroom requirements, showing interest in pursuing knowledge independently	2 pts Poor Relies on classroom instruction only	1 pts Very Poor No initiative or interest in acquiring new knowledge	6 pts
⌚ SO 7 PI 2 Student Outcome 7.2 Learn independently threshold: 4.8 pts	6 pts Excellent Completes an assigned task independently and practices continuous improvement	5 pts Good Completes an assigned task without supervision or guidance	4 pts Satisfactory Requires minimal guidance to complete an assigned task	3 pts Unsatisfactory Requires detailed or step-by-step instructions to complete a task	2 pts Poor Shows little interest to complete a task independently	1 pts Very Poor No interest to complete a task independently	6 pts
⌚ SO 7 PI 3 Student Outcome 7.3 Critical thinking in the broadest context of technological change threshold: 4.8 pts	6 pts Excellent Synthesizes and integrates information from a variety of sources; formulates a clear and precise perspective; draws appropriate conclusions	5 pts Good Evaluate information from a variety of sources; formulates a clear and precise perspective.	4 pts Satisfactory Analyze information from a variety of sources; formulates a clear and precise perspective.	3 pts Unsatisfactory Apply the gathered information to formulate the problem	2 pts Poor Gather and summarized the information from a variety of sources but failed to formulate the problem	1 pts Very Poor Gather information from a variety of sources	6 pts
⌚ SO 7 PI 4 Student Outcome 7.4 Creativity and adaptability to new and emerging technologies threshold: 4.8 pts	6 pts Excellent Ideas are combined in original and creative ways in line with the new and emerging technology trends to solve a problem or address an issue.	5 pts Good Ideas are creative and adapt the new knowledge to solve a problem or address an issue	4 pts Satisfactory Ideas are creative in solving a problem, or address an issue	3 pts Unsatisfactory Shows some creative ways to solve the problem	2 pts Poor Shows initiative and attempt to develop creative ideas to solve the problem	1 pts Very Poor Ideas are copied or restated from the sources consulted	6 pts